

# A Brief Introduction to Parallel Computing

Zhang, Honglin

Department of Physics, Tsinghua University


June 25, 2012

Prologue

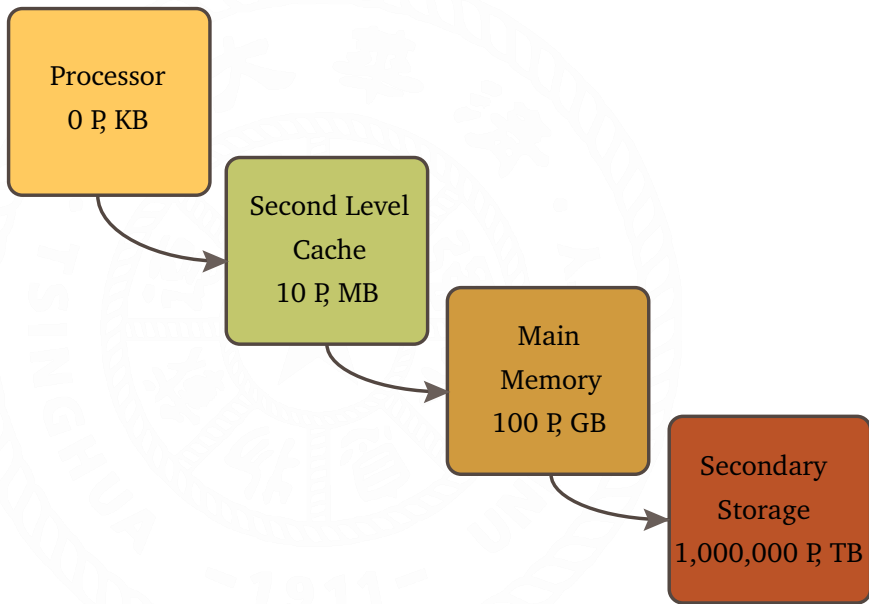
# Why Parallel Computing?

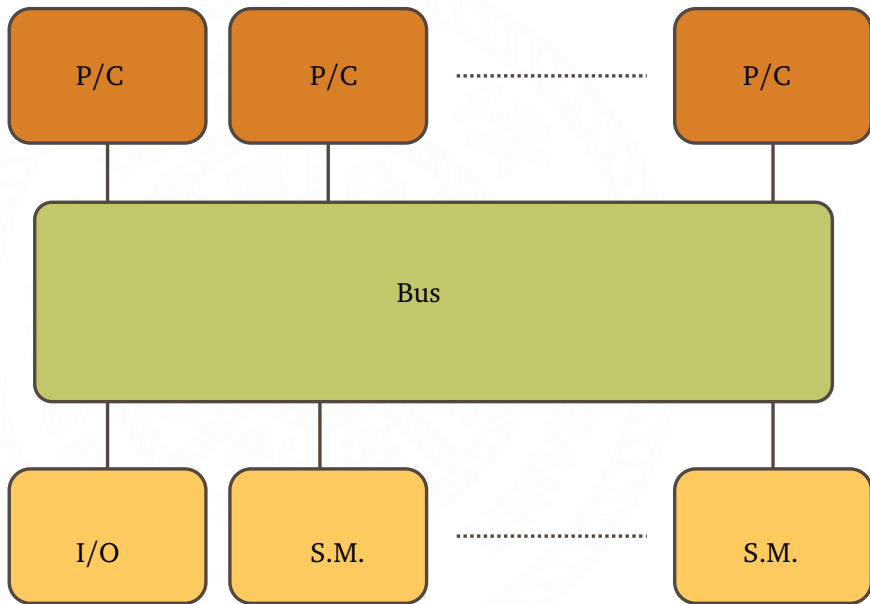
Quick Review


# Introduction to Hardware



Device	Pro.	Sec. Lev.	Cac.	Main Mem.	Sec. Sto.
CPU Period	0		10	100	1,000,000
Speed	10GB/s		2GB/s	1GB/s	10MB/s
Size	KB		MB	GB	TB









Before we hit the road,

# Processes vs. Threads

## Processes

Resource Holder

Independent Execution

Programs Loaded as

Processes by OS

## Processes

Resource Holder  
Independent Execution  
Programs Loaded as  
Processes by OS

## Threads

(Master) Thread at Program Entry Point  
More Threads Created by Existing Ones Possible  
Separate Stack for each thread  
Shared Code and Data Segments

# Why Threads?

- Efficient Data Sharing

# Why Threads?

- Efficient Data Sharing
- High Multi-core Performance

Typical Parallel Models I.

# Shared Memory Model



# Shared Memory Model

- Unified Address Space

# Shared Memory Model

- Unified Address Space
- Asynchronous Execution



# Shared Memory Model

- Unified Address Space
- Asynchronous Execution
- Able to Explicitly Synchronize

# Shared Memory Model

- Unified Address Space
- Asynchronous Execution
- Able to Explicitly Synchronize
- Local Variables along with Shared Variables

# Shared Memory Model

- Unified Address Space
- Asynchronous Execution
- Able to Explicitly Synchronize
- Local Variables along with Shared Variables
- Threads Created and Destroyed Dynamically

# A Typical Terrible Example

```
static int s=0;
```

## Thread 1

```
reg0=f(A[i]);  
reg1=s;  
reg1=reg1+reg0;  
s=reg1;
```

## Thread 2

```
reg0=f(A[i]);  
reg1=s;  
reg1=reg1+reg0;  
s=reg1;
```

# A Typical Terrible Example

```
static int s=0;
```

## Thread 1

```
reg0=f(A[i]);  
reg1=s;  
reg1=reg1+reg0;  
s=reg1;
```

## Thread 2

```
reg0=f(A[i]);  
reg1=s;  
reg1=reg1+reg0;  
s=reg1;
```

Wait a moment

# What value will “s” take?

Eureka!

Here is the problem! There are more than one threads visiting the variable “s” at the same moment.

Eureka!

Here is the problem! There are more than one threads visiting the variable “s” at the same moment.

Remember

This is very important!

# A Typical Good Example

```
static int s=0;  
static lock lk;
```

## Thread 1

```
local_s1=0;  
for i=0, n/2-1  
    local_s1=local_s1+f(A[i]);  
lock(lk);  
s=s+local_s1;  
unlock(lk);
```

## Thread 2

```
local_s2=0;  
for i=0, n/2-1  
    local_s2=local_s2+f(A[i]);  
lock(lk);  
s=s+local_s2;  
unlock(lk);
```



To implement shared memory model, we usually adopt OpenMP.

To implement shared memory model, we usually adopt OpenMP.

## Tips

OpenMP is not a language,  
but a compiler directive!

To implement shared memory model, we usually adopt OpenMP.

## Tips

OpenMP is not a language,  
but a compiler directive!

Supported by:

- C/C++/Fortran
- Sun Compiler/GNU Compiler/Intel Compiler

Typical Parallel Models II.

# Message Passing Model



# Message Passing Model

- Separate Address Space

# Message Passing Model

- Separate Address Space
- Explicit Send/Receive Pair

# Message Passing Model

- Separate Address Space
- Explicit Send/Receive Pair
- Explicitly Synchronize as Default

# Message Passing Model

- Separate Address Space
- Explicit Send/Receive Pair
- Explicitly Synchronize as Default
- Process Number Decided Before Launch



Message passing involves message sending and receiving.

Message passing involves message sending and receiving.

Two typical forms of communication:

- Blocking, e.g. **Telephone**

Message passing involves message sending and receiving.

Two typical forms of communication:

- Blocking, e.g. **Telephone**
- Non-blocking, e.g. **Post Office**

Emm,

# Feel the difference?

# A Typical Terrible Example

## Process 1

```
xlocal=A[1];  
send xlocal, proc2;  
receive xremote, proc2;  
s=xlocal+xremote
```

## Process 2

```
xlocal=A[2];  
send xlocal, proc1;  
receive xremote, proc1;  
s=xlocal+xremote
```

# A Typical Terrible Example

## Process 1

```
xlocal=A[1];  
send xlocal, proc2;  
receive xremote, proc2;  
s=xlocal+xremote
```

## Process 2

```
xlocal=A[2];  
send xlocal, proc1;  
receive xremote, proc1;  
s=xlocal+xremote
```

Wait a moment

# Will there be an interlocking?

Eureka!

Send & Receive are blocking operation. The sender will wait silently until the message is received by a receiver.

Eureka!

Send & Receive are blocking operation. The sender will wait silently until the message is received by a receiver.

Remember

This is very important!

# A Typical Good Example

## Process 1

```
xlocal=A[1];  
send xlocal, proc2;  
receive xremote, proc2;  
s=xlocal+xremote;
```

## Process 2

```
xlocal=A[2];  
receive xremote, proc1;  
send xlocal, proc1;  
s=xlocal+xremote;
```



# A Typical Good Example

## Process 1

```
xlocal=A[1];  
send xlocal, proc2;  
receive xremote, proc2;  
s=xlocal+xremote;
```

## Process 2

```
xlocal=A[2];  
receive xremote, proc1;  
send xlocal, proc1;  
s=xlocal+xremote;
```

Of course, we may also implement this in a non-blocking way...

# A Typical Good Example

## Process 1

```
xlocal=A[1];  
send xlocal, proc2;  
receive xremote, proc2;  
s=xlocal+xremote;
```

## Process 2

```
xlocal=A[2];  
receive xremote, proc1;  
send xlocal, proc1;  
s=xlocal+xremote;
```

Of course, we may also implement this in a non-blocking way...

However, non-blocking is not always reliable.

To implement message passing model, we usually adopt MPI.

To implement message passing model, we usually adopt MPI.

## Tips

**MPI is a library, not a language!**

To implement message passing model, we usually adopt MPI.

## Tips

**MPI is a library, not a language!**

Supporting:

- C/C++/Fortran

Want to know more?

# More About OpenMP

Obtain the number of threads:

C/C++: `int omp_get_num_threads(void)`

Fortran: `integer function omp_get_num_threads()`

## Obtain the number of threads:

C/C++: `int omp_get_num_threads(void)`

Fortran: `integer function omp_get_num_threads()`

## Obtain the thread number

C/C++: `int omp_get_thread_num(void)`

Fortran: `integer function omp_get_thread_num()`

# Notice thread 0 is the master!



## Obtain the number of processors

C/C++: `int omp_get_num_procs(void)`

Fortran: `integer function omp_get_num_procs()`

## Obtain the number of processors

C/C++: `int omp_get_num_procs(void)`

Fortran: integer function `omp_get_num_procs()`

## Set the number of threads

C/C++: `void omp_set_num_threads(int number_threads)`

Fortran: subroutine `omp_set_num_threads(number_threads)`

More for you to find

<https://computing.llnl.gov/tutorials/openMP/>

<http://openmp.org/wp/openmp-specifications/>

Still More?

# More About MPI

MPI defines a library for C, Fortran and Java. MPI's definition includes no specific features of any manufacturer, OS or device.

MPI defines a library for C, Fortran and Java. MPI's definition includes no specific features of any manufacturer, OS or device.

Especially, MPICH environment is broadly accepted.

## MPICH Download

<http://www.mcs.anl.gov/research/projects/mpich2>

## Compile in MPICH

- mpicxx/mpicc/mpif77/mpif90
- -cc=icc/gcc designates a specific compiler

## Launch a MPI program

```
mpiexec -n N routine
```

```
mpiexec -host HOST_1 -n N routine_1 : -host HOST_2 -n M routine_2
```



## Message Envelope for Sender

`MPI_Send(address,count,datatype,destination,tag,communicator)`

### Description

address: data pointer

count: length of datatype

datatype: unified data type in MPI

destination: ID of process receiver

tag: for programmer use

communicator: ID of communication group

## Message Envelope for Receiver

`MPI_Recv(address,count,datatype,source,tag,communicator,status)`

## Description

status: a pointer of datatype `MPI_Status`, which record the information of sending and receiving

## Broadcast

`MPI_Bcast(address,count,datatype,root,comm)`

## Scatter

For Root:

```
MPI_Scatter(send_address,send_count,send_datatype,root,comm)
```

For others:

```
MPI_Scatter(recv_address,recv_count,recv_datatype)
```

## Gather

For Root:

```
MPI_Gather(recv_address,recv_count,recv_datatype,root,comm)
```

For others:

```
MPI_Gather(send_address,send_count,send_datatype)
```

## Barrier

**MPI\_Barrier(comm)**

## Description

Let processes in the same communicator wait for each other.

More for you to find

<http://www.mcs.anl.gov/research/projects/mpi/>

Not Finished Yet?

# About GPU Programming



So far, we have not mentioned CUDA programming yet.

CUDA is so *unfriendly* that I would not address here.



More for you to find

[http://www.nvidia.com/object/cuda\\_home\\_new.html](http://www.nvidia.com/object/cuda_home_new.html)

## Advanced Tips for HPC?

- Profiling tools, e.g. VTune.

## Advanced Tips for HPC?

- Profiling tools, e.g. VTune.
- Find senior programmers for help.

## Advanced Tips for HPC?

- Profiling tools, e.g. VTune.
- Find senior programmers for help.
- Parallelize your program only when you have to.

## Advanced Tips for HPC?

- Profiling tools, e.g. VTune.
- Find senior programmers for help.
- Parallelize your program only when you have to.



Q&A

# Any Question?

Epilogue

# Thanks for your Attention!